

CPSC 213 – Assignment 2

CPU and Static Scalars and Arrays

Due: Monday, Oct 3, 2016 at 11:59pm
No late assignments accepted.

Goal

In this assignment you will implement a significant subset of the SM213 ISA we are developing in class: the memory-access and arithmetic instructions that are listed below.

You will then explore how these instructions are used to implement global scalars and arrays (both static and dynamic) in C by carefully examining code snippets in Java, C and assembly language. An important part of this evaluation is to carefully observe the dynamic behaviour of the assembly-language snippets by executing them in the simulator. You will develop intuition about the connection between the high-level language statements, their machine-code implementation and the execution of this code by the CPU hardware.

By implementing these instructions in the simulator you will see what is required to build them in hardware and you will deepen your understanding of what a global variable is, what memory is, and what role that compiler and hardware play in implementing them.

A key thing to think about while doing this is: “what does the compiler know about these variables” and so what can the compiler hard-code in the machine code it generates. For global variables, recall that the compiler knows their address. And so the address of a global variable is hardcoded in the instructions that access it. But, the compiler does not know the address of a dynamic array and so, even though it knows the address of the variable that stores the array reference, it must generate code to read the array’s address from memory when the program runs.

Implement these SM213 Instructions

Implement the following instructions in the SM213 Simple Machine Simulator by modifying the `execute()` method of the CPU class. You will also need a completed implementation of the `MainMemory` class. You can use your implementation from Assignment 1 or the solution, which is available on the course web page.

The instructions are described in more detail, including examples, in the *213 Companion*.

Many of these instructions, but not all, are used in this week’s code snippets.

Instruction	Assembly	Format	Semantics
load immediate	ld \$v, rd	0d-- vvvvvvvv	$r[d] \leftarrow v$
load base + offset	ld o(rs), rd	1psd	$r[d] \leftarrow m[o=p*4+r[s]]$
load indexed	ld (rs,ri,4), rd	2sid	$r[d] \leftarrow m[r[s]+r[i]*4]$
store base + offset	st rs, o(rd)	3spd	$m[o=p*4+r[d]] \leftarrow r[s]$
store indexed	st rs, (rd,ri,4)	4sdi	$m[r[d]+r[i]*4] \leftarrow r[s]$

Memory-Access Instructions (and load immediate)

Instruction	Assembly	Format	Semantics
rr move	mov rs, rd	60sd	$r[d] \leftarrow r[s]$
add	add rs, rd	61sd	$r[d] \leftarrow r[d] + r[s]$
and	and rs, rd	62sd	$r[d] \leftarrow r[d] \& r[s]$
inc	inc rd	63-d	$r[d] \leftarrow r[d] + 1$
inc addr	inca rd	64-d	$r[d] \leftarrow r[d] + 4$
dec	dec rd	65-d	$r[d] \leftarrow r[d] - 1$
dec addr	deca rd	66-d	$r[d] \leftarrow r[d] - 4$
not	not rd	67-d	$r[d] \leftarrow \sim r[d]$
shift	shl \$v, rd shr \$v, rd	7dss	$r[d] \leftarrow r[d] \ll s$ <i>s = v for left and -v for right</i>
halt	halt	f000	throw halt exception
nop	nop	ff00	do nothing (nop)

ALU Instructions

Code Snippets Used this Week

The file [a2_code.zip](#) contains code snippets for Part I as well as other files for Part II, described below.

You will use the following code snippets this week. There are C, Java and SM213 Assembly versions of each of these (except the C-pointer-math file, for which there is no Java).

- S1-global-static
- S2-global-dyn-array

Your Instruction Implementation [70%]

Implement the `execute()` method of the CPU class in the `arch.sm213.machine.student` package. This method uses the register values stored by the fetch stage to execute the instructions, transforming the register file (i.e., `reg`) and main memory (i.e., `mem`) appropriately. The changes you need to make are indicated by “TODO” flags.

Testing and Debugging Your Implementation [30%]

The simulator displays the current value of the register file, main memory and the internal registers such as the pc and instruction registers. Use the simulator to test and debug. If necessary, you can also set breakpoints in your CPU class, but most of the debugging can probably be done without doing this just by examining how the machine state changes and by paying careful attention to exception messages the simulator displays at the bottom of the window.

The first thing you will want to do is to create a simple test assembly file with various forms of the instructions that you implement. Name this file “`test.s`”. Use this file to test and debug each instruction in turn. To test an instruction, write a couple of lines of assembly that use that instruction in `test.s` and then observe what the simulator does when you step through these instructions. For example, to test the *load-immediate* instruction, you might add lines like these to `test.s`:

```
ldi:  ld $0x11223344, r0
      ld $0x11223344, r7
      halt
```

Load `test.s` into the simulator and use its GUI to set initial values for `r0` and `r7`, to step through these instructions, and to verify that the `r0`’s and `r7`’s ending values are `0x11223344` and that the values of the other registers did not change.

An alternative and optional testing approach is explained at the end of this document.

Suggested Implementation Approach

The most important aspect of any strategy for implementing complicated software is to test as you go. Applied in this context, this might mean implementing each instruction in turn, testing each one before moving to the implementation of the next.

You will likely find that implementing and debugging the first instruction is the hardest. Once you have one working, you will see that adding others will follow a pattern.

Once you've tested every instruction, then run the snippets to observe what they do.

Using the Simulator

You'll get help using the simulator in the labs, but here are a few quick things that you will find helpful.

1. You can edit instructions and data values directly in the simulator (including adding new lines or deleting them).
2. The simulator allows you to place "labels" on code and data lines. This label can then be used as a substitute for the address of those lines. For example, the variable's `a` and `b` are at addresses `0x1000` and `0x2000` respectively, but can just be referred to using the labels `a` and `b`. Keep in mind, however, that this is just an simulator/assembly-code trick, the machine instructions still have the address hardcoded in them. You can see the machine code of each instruction to the left of the instructions in the *memory image* portion of the instruction pane.
3. You can change the program counter (i.e., `pc`) value by double-clicking on an instruction. And so, if you want to execute a particular instruction, double click it and then press the "Step" button. The instruction pointed to by the `pc` is coloured green.
4. Memory locations and registers read by the execution of an instruction are coloured blue and those written are coloured red. With each step of the machine the colours from previous steps fade so that you can see locations read/written by the past few instructions while distinguishing the cycle in which they were accessed.
5. Instruction execution can be animated by clicking on the "Show Animation" button and then single stepping or running slowing.

Executing the Snippets

Once you have your implementation of the Simulator for these instructions working, the fun has only just begun. You can now execute this week's two snippets in the simulator to see what happens when they run. Single step through each of the snippets and observe what changes in the register file and/or main memory as the result of each instruction. There is nothing to turn in for this step. Just used these examples and the simulator to help you understand the code.

***NOTE** that you have also been provided a reference implementation of the simulator. Assignment 1 explains how to run it. You can use this implementation along side of your implementation, to verify its correctness, or instead of your implementation if you haven't finished it by the time you want to do the rest of the assignment.*

What to Hand In

Use the `handin` program to hand in the following in a directory named `a2`. Please do not hand in anything that isn't listed below. In particular, **do not hand in your entire Eclipse workspace nor the entire source tree for the simulator and do not hand in any class files.**

1. `README.txt` – that contains the name and student number of you and your partner
2. `PARTNER.txt` – containing your partner's CS login id and nothing else (i.e., the 4- or 5-digit id in the form a0z1). Your partner should not submit anything.
3. `CPU.java` – that implements the instructions listed above
4. `test.s` – the assembly program you used to test the instructions.
5. `TEST.txt` – a plain-text file (not word or pdf) that describes your test procedure

OPTIONAL – Testing Using the Command Line

The simulator has a command-line (i.e., non-GUI) interface. This would allow you, for example, to build a test script to automate regression testing.

To invoke the command-line version of the simulator you need to locate three items in your file system: your implementations of `MainMemory.class` and `CPU.class` and `SimpleMachineStudent213.jar` that you downloaded for this assignment (or Assignment 1). Lets say that all three are in the same directory and you have “**cd**”ed into that directory (i.e., if you type `ls` you see these three files) and thus the directory is called “.”. To invoke the command-line simulator you type:

```
java -cp .:SimpleMachineStudent213jar SimpleMachine -i cli -a sm213 -v student
```

Then type `help` to see a list of commands. Note that register names are distinguished from label names by adding a percent sign to the register name; e.g., `%r0` is the name for register 0.

To run test the load instruction you might type the following (what you type is in bold):

```
% java -cp .:SimpleMachineStudent213jar SimpleMachine -i cli -a sm213 -v student
Simple Machine (SM213-Student)
(sm) l test.s
(sm) %r0=0
(sm) %r7=0
(sm) s
(sm) s
(sm) i reg
%r0: 0x11223344 287454020
%r1: 0x0        0
%r2: 0x0        0
%r3: 0x0        0
%r4: 0x0        0
%r5: 0x0        0
%r6: 0x0        0
%r7: 0x11223344 287454020
```

If you want to run it again, or when you have several tests and you want to go to a specific one, you can use the `goto` command with the label associated with the first instruction of that test. And, if you place a `halt` instruction after each test, you can use the `run` command instead of stepping. For example, in this case:

```
(sm) g ldi
(sm) r
```

Note that if you step past the end of the the loaded file you will get an address out of bounds exception.

Note also that if you want to run the reference implementation from the command line you leave off the “-a” and “-v” command line flags and type the following instead.

```
java -jar SimpleMachine213.jar -i cli
```

You can use the **assert** command to streamline the testing to look something like this.

```
% java -cp .:SimpleMachineStudent213jar SimpleMachine -i cli -a sm213 -v student
```

Simple Machine (SM213-Student)

```
(sm) l test.s
(sm) %r0=0
(sm) %r7=0
(sm) g ldi
(sm) r
(sm) a "ld imm"
(sm) a %r0==0x11223344
(sm) a %r1==0
(sm) a %r7==0x11223344
```

If an assertion succeeds, then it prints nothing. If it fails it prints something like this:

```
XXX ASSERTION FAILURE (ld imm): %r0 == 0x11223344 != 0x0
```

You could use this approach to run the entire test in a script and then scan its output for lines that indicate an assertion failure.

The script is simply a file that contains the commands to execute in sequence, for example, the file named **test-script** might look like this:

```
l test.s
%r0=0
%r1=0
g ldi
r
a "ld imm"
a %r0==0x11223344
a %r1==0
a %r2==0x11223344
```

You can then use the unix shell "<" operator to run the simulator with this as input.

```
java -cp... -v student < test-script
```

And you can then process the output to look for assertion failures by using a similar trick to pipe the output to a unix command called **grep** using the "|" operator like this.

```
java -cp... -v student < test-script | grep XXX
```

The resulting output is a list of the tests that failed.

All of this is entirely optional and it will take some work to setup, so don't worry about trying this if you aren't up for the challenge. We'll help you if you do want to try. The advantage of investing time on the setup is that testing will run a bit more smoothly for you.